

**Faculty of Engineering & Technology**  
**Electrical & Computer Engineering Department**  
**Operating System Concepts – ENCS3390**

**Project #1**

**Multi- Processing and Multi- Threading Programming**

---

**Prepared by:** Khaled Abu Lebdeh

**ID Number:** 1220187

**Instructor:** Dr. Yazan Abu Farha

**Section:** 1

**Date:** November 25, 2024

## Abstract

Programmers used to write sequential codes. However, large or heavy programs may take long execution time even if they implemented in optimal algorithms.

Sequential code does not take the full advantage of multi- core systems. On the contrary, there are different approaches use multi- core systems to increase the total throughput, like Multi- Processing and Multi- Threading.

## Table of Contents

### Contents

<b>Abstract .....</b>	<b>ii</b>
<b>Procedure.....</b>	<b>4</b>
1- Computer Environment.....	4
2- Multiprocessing and Multithreading Requirements.....	4
2-1 Multiprocessing Achievement.....	4
2-2 Multithreading Achievement .....	6
3- Results and Performance of approaches used.....	8
4- Analysis using Amdahl's Law .....	14
5- Comments on different approaches .....	16
<b>Conclusion.....</b>	<b>17</b>

## Procedure

### 1- Computer Environment

- The computer used to execute programs has (intel core i5, 8<sup>th</sup> Gen) processor. Processor's base (guaranteed) speed is 1.8 GHz. However, the average speed while execution was about 3.2 GHz.
- A virtual machine with “*Linux System*” was used with 5-cores and 4 GB memory.
- Programs were written in “*C- Language*” and ran within Linux terminal.


### 2- Multiprocessing and Multithreading Requirements

#### 2-1 Multiprocessing Achievement

The algorithm used was to load words into an array of strings. Then, processes were created using **fork()** API, each process was given a unique number named as “process\_num” which determines start and end indices of words array the process must work on.

$$\text{start} = (\text{number of array cells}) * (\text{process\_num}) / (\text{number of total processes created})$$

$$\text{end} = (\text{number of array cells}) * (\text{process\_num} + 1) / (\text{number of total processes created})$$



A shared memory was created using **shm\_open()** and **mmap()** APIs to allow processes to communicate and store results gotten from all processes. Shared memory took form of a two- dimensional array, each process found the

frequency of each word in its data slot, and stored the result in the shared memory as a struct (a word and its frequency). Each process wrote on one row of the 2-D array only, to avoid race condition.


Children processes used **exit()** API after finishing execution, to terminate. While children processes hadn't finished executing, the parent process was waiting for them using **wait()** API. After that, the parent merges children's results using a *Binary Search Tree* (BST) data structure to find total frequency of each word.

Parent process used Min-Heap data structure to find top 10 frequent words in the BST containing all words with their frequencies. For each tree node, we enqueued the word into the heap if its frequency is larger than min frequency in heap. So, the heap then contained the top 10 frequent words among all words!

APIs used in Multiprocessing	Description
fork()	Creates child process.
exit()	Process uses it to terminate.
wait()	Parent process uses it to wait for children processes.
shm_open() and mmap()	Creates shared memory.
memset()	To set initial values, it was used in the code to reset shared data to zeros.
malloc()	Allocates memory during runtime.
free()	Deallocates memory allocated during runtime.

## 2-2 Multithreading Achievement

The algorithm used was to load words into an array of strings. All variables and arrays needed in threads were declared in the global scope. Then, an array of *pthread\_t* type (from *pthread.h* library) was declared to hold threads identifiers (thread\_ID). Threads were created in a loop using **pthread\_create()** API and given a unique number “thread\_num” which determines start and end indices of words array the thread must work on. Another loop was needed to wait for threads’ termination using **pthread\_join()** API.



Each thread executed a function usually named as `runner()`, passed as an argument through **pthread\_create()** API. The `runner()` functions called user-defined functions needed to count words within thread’s data slot. Runner function initialized temporary array to hold counted frequency words in its data slot, to avoid writing on global data and causing race condition. Binary Search Algorithm was used for optimal words search. After all threads had terminated using **pthread\_exit()** API, the total frequency was merged into the global data array using BS. Then a Min-Heap was used to find top 10 frequent words among the merged array (similar to one in Multiprocessing).

- AI was used to add BS algorithm to the code

APIs used in Multithreading	Description
pthread_create()	Creates thread.
pthread_exit ()	thread uses it to terminate.
pthread_join()	Wait for threads to terminate.
malloc()	Allocates memory during runtime.
free()	Deallocates memory allocated during runtime.

### 3- Results and Performance of approaches used

The task is to find top 10 frequent words in enwik8 dataset. We found results using three different approaches, Naive approach (a program that does not use any child processes or threads), Multiprocessing approach (a program that uses multiple child processes running in parallel) and Multithreading approach (a program that uses multiple joinable threads running in parallel).

Output:

```
1 | the      -- 1060930 times
2 | of       -- 593543 times
3 | and      -- 416437 times
4 | one      -- 411559 times
5 | in       -- 372213 times
6 | a        -- 326320 times
7 | to       -- 316319 times
8 | zero     -- 264789 times
9 | nine     -- 250239 times
10| two      -- 192535 times
```

Approach	Execution Time
Naïve	45 minutes
Multi- processing (2 processes)	27 minutes
Multi- processing (4 processes)	16.19 minutes
Multi- processing (6 processes)	15.1 minutes
Multi- processing (8 processes)	12.2 minutes
Multi- threading (2 threads)	28 seconds
Multi- threading (4 threads)	16.7 seconds
Multi- threading (6 threads)	16.3 seconds
Multi- threading (8 threads)	15.2seconds



```

file has been read successfully.
Find top 10 frequent words in Naive approach
two      --      192644 times
nine     --      250430 times
a        --      325873 times
and      --      416629 times
zero     --      264975 times
of       --      593677 times
the      --      1061396 times
in       --      372201 times
one      --      411764 times

Dynamic allocated memory have been freed

real    45m12.459s
user    45m5.455s
sys      0m6.460s
File has been read successfully

```

*Figure 1 : Naïve Approach's output*

```

File has been read successfully

Enter number of processes you want to run: 2
zero      --      141382 times
nine      --      143201 times
a         --      173705 times
in        --      196333 times
to        --      165725 times
and       --      219230 times
are       --      39828 times
of        --      311482 times
the       --      551993 times

real    27m40.915s
user    54m29.324s

```

*Figure 2 : Multiprocessing Approach's output (2 processes)*

```

File has been read successfully

Enter number of processes you want to run: 4
is          --          4855 times
zero        --          6432 times
a           --          8489 times
in          --          9922 times
one         --          8243 times
and         --          11692 times
as          --          3667 times
of          --          15729 times
the         --          27688 times

real        16m19.078s
user        59m44.622s
sys         0m29.863s

```

*Figure 3 : Multiprocessing Approach's output (4 processes)*

```

File has been read successfully

Enter number of processes you want to run: 6
nine        --          250430 times
zero        --          264975 times
a           --          325873 times
in          --          372201 times
to          --          316376 times
and         --          416629 times
as          --          131815 times
of          --          593677 times
the         --          1061396 times

real        15m1.219s
user        77m14.512s
sys         1m29.213s

```

*Figure 4 : Multiprocessing Approach's output (6 processes)*

```

File has been read successfully

Enter number of processes you want to run: 8
nine      --      250430 times
zero      --      264975 times
a         --      325873 times
in        --      372201 times
to        --      316376 times
and       --      416629 times
as        --      131815 times
of        --      593677 times
the       --      1061396 times

Process returned 0 (0x0)   execution time : 734.669 s
Press ENTER to continue.

```

*Figure 5 : Multiprocessing Approach's output (8 processes)*

```

File has been read successfully

Enter number of threads: 2
Number of threads = 2

Top 10 frequent words:
the       --      1060930 times
of        --      593543 times
and       --      416437 times
one       --      411559 times
in        --      372213 times
a         --      326320 times
to        --      316319 times
zero      --      264789 times
nine      --      250239 times
two       --      192535 times

Dynamically allocated memory has been freed

real      0m27.908s
user      0m50.843s
sys       0m0.503s

```

*Figure 6 : Multithreading Approach's output (2 threads)*

```

File has been read successfully

Enter number of threads: 4
Number of threads = 4

Top 10 frequent words:
the      --      1060930 times
of       --      593543 times
and      --      416437 times
one      --      411559 times
in       --      372213 times
a        --      326320 times
to       --      316319 times
zero     --      264789 times
nine     --      250239 times
two      --      192535 times

Dynamically allocated memory has been freed

real    0m16.766s
user    0m53.770s
sys     0m0.391s

```

Figure 7 : Multithreading Approach's output (4 threads)

```

6File has been read successfully

Enter number of threads:
Number of threads = 6

Top 10 frequent words:
the      --      1060930 times
of       --      593543 times
and      --      416437 times
one      --      411559 times
in       --      372213 times
a        --      326320 times
to       --      316319 times
zero     --      264789 times
nine     --      250239 times
two      --      192535 times

Dynamically allocated memory has been freed

real    0m16.381s
user    1m17.500s
sys     0m0.505s

```

Figure 8 : Multithreading Approach's output (6 threads)

```
8File has been read successfully

Enter number of threads:
Number of threads = 8

Top 10 frequent words:
the      --      1060930 times
of       --      593543 times
and      --      416437 times
one      --      411559 times
in       --      372213 times
a        --      326320 times
to       --      316319 times
zero     --      264789 times
nine     --      250239 times
two      --      192535 times

Dynamically allocated memory has been freed

real    0m15.292s
user    1m5.162s
sys     0m0.557s
```

*Figure 9 : Multithreading Approach's output (8 threads)*

#### 4- Analysis using Amdahl's Law

Amdahl's Law identifies performance gains from adding additional cores to an application that has both serial and parallel components.

$$\text{Speedup} \leq \frac{1}{S + \frac{1-S}{N}}$$

S: is serial portion of code.

P: is parallel portion of code.

N: number of cores

$P = (1-S)$ .

- To find the series portion of code, we measure time before and after parallel section of code, then divide by the total time to get the series portion.

We found that **S=12%** approximately. So, **P= 100% - S%= 88%**

The reason for the high portion of parallel code is that complex operations are done in parallel. Only few tasks were implemented in series, like reading the file and printing the top 10 frequent words. The more parallel portion leads to more efficiency on multi-core systems.

- The max speedup according to the available number of cores occurs in ideal scenarios if  $S=0\%$  and  $P=100\%$

**Speedup  $\leq N$  times, which means Speedup  $\leq 5$  times**

- The optimal number of threads or processes when working on a system with 5 cores is actually about (4 or 5) threads or processes. The reason is that no more than 5 threads or processes can work in parallel, so they allocate resources with no worth improvements.

## 5- Comments on different approaches

- In Naïve approach, the execution time was relatively high, since there is no more than one thread of execution running at a time.
- In Multiprocessing approach, starting with two parallel processes, the improvement and difference in terms of execution time was very clear. The same scenario occurred with four parallel processes. However, increasing in terms of number of processes does not guarantee much improvement in execution time, since there are only 5 cores allocated to virtual machine. If there are more than five processes, they can't all run in parallel. The system can run five parallel processes at max. That resulted in no much improvement when six or eight processes were used.
- In multithreading approach, the execution time was incredible improved than both Naïve and Multiprocessing approaches. The reason is that Multithreading supports parallel execution with light- weight processes. Thread creation is much saver than process creation. Moreover, thread switching is lower overhead than process context switching. In addition, threads share code section and global scope section, so no shared memory needed to communicate between threads. However, as same as in multiprocessing, increasing in terms of number of threads does not guarantee much improvement in execution time. It depends on the system used.

In general, the relation between multiple threads of execution and speedup isn't linear. It is somehow an **exponential** relation. (assume code with parallel portion).



## Conclusion

Today's computers are interactive ones, dealing with user interactions and different I/Os. That led to the motivation of throughput improvement. Many tasks are real-time ones, which must be executed in specific turnaround time.

Completely sequential programs (with 0% parallel portion) probably do not get the advantage of today's multi-core computers. So, it is somehow required to move towards parallel execution.

Parallel execution could be implemented by multiple processes or threads. If the program can be implemented using threads, it is much better than parallel processes. Parallel execution often satisfies users required throughput. However, explicit parallel coding could be difficult for programmers. Programmers should be responsible for dividing activities, load balancing and data dependency. That was a motivation to use implicit threading!